Module 6 Logistic Regression and Applying GLM

August 11, 2024

0.1 Dempsey Wade

1 Module 6: Logistic Regression and Applying GLM

Authors: Favio Vázquez and Jessica Cervi

In this assignment, you will be working on more complicated Python functions, looking at Logistic Regression Models and how to apply GLM to them.

We begin by importing most of the libraries you will need throughout this assignment.

Index:

- Question 1
- Question 2
- Question 3
- Question 4
- Question 5
- Question 6
- Question 7
- Question 8

1.1 Import Libraries

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

1.2 Python Functions

Return to top

1.3 Question 1

Create a Python function called multiply that takes a variable number of arguments and multiplies them. Then multiply the numbers 1,6,3 and 8 and assign the result to a variable called prod_total.

```
[2]: ### GRADED
     ### YOUR SOLUTION HERE
     def multiply(*args):
         total = 1
         for i in args:
             total = total*i
         return total
     prod_total = multiply(1,6,3,8)
     prod_temp = multiply(5,5)
     prod_temp = multiply(8, 7)
     print(prod_temp)
     ###
     ### YOUR CODE HERE
     ###
     ### Answer check
     print("Result: {}".format(prod_total))
```

56 Result: 144

```
[3]: ###
### AUTOGRADER TEST - DO NOT REMOVE
###
```

Return to top ## Question 2

Python functions can return multiple values. Create a Python function called calculator that takes two integers a and b and returns a tuple. The first element of the tuple must be the sum of a and b. The second element must be the square of their product.

Call the function with a=6, b=10 and assign the sum to add_values and the squared product to square_values, respectively.

```
[4]: ### GRADED

### YOUR SOLUTION HERE
a = 6
b = 10

def calculator(a, b):
   add = a+b
   square = (a*b)*(a*b)
   return add, square

add_values, square_values = calculator(a,b)
```

```
###
### YOUR CODE HERE
###
### Answer check
print("Sum: {}\nSquared product: {}".format(add_values, square_values))
Sum: 16
```

Squared product: 3600

Return to top ## Question 3

Define a function called **subtract** that takes two integers and computes the absolute value of their difference. The result must not be returned but set to the global variable **difference** instead.

Next, call the function subtract with arguments 7 and 9 and check the value of difference.

Hint: Do not assign the call of **subtract** to **difference** since the function returns **None** and doing so will override the current value of the global variable.

```
[6]: ### GRADED
import math

### YOUR SOLUTION HERE
a = 7
b = 9

def subtract(a, b):
    global difference
    difference = (abs(a-b))

subtract(a,b)
print(difference)

###
### YOUR CODE HERE
###

### Answer check
print("The value of difference is: {}".format(difference))
```

2
The value of difference is: 2

```
[7]: ###
### AUTOGRADER TEST - DO NOT REMOVE
###
```

1.4 Logistic Regression Models

In the first part, we will use a fake Instagram data set, indicating whether or not a particular instagram user clicked on an ad in their feed. We will create a model that predicts the likelihood of clicking on another ad based off the features of that user.

The data set contains the following features:

- 'Daily Time Spent on Instagram': user time spent on app in minutes
- 'Age': user age in years
- 'Daily Internet Usage': Avg. minutes a day consumer is on the internet
- 'Ad feed headline': Headline of the Ad
- 'Location': current location of consumer
- 'Sex': 1 for male and 0 for female
- 'Country': Country of consumer
- 'Timestamp': Time at which consumer clicked on Ad or closed window
- 'Clicked on Feed': 0 or 1 indicated clicking on Feed

1.4.1 Read the Data

 $insta_data = pd.read_csv('data/insta.csv')$

```
[8]: #Read in data
insta_data = pd.read_csv('/Users/dempseywade/Desktop/gitRepo/

→DartmouthCodingAssignments/data/Mod6_insta.csv')
```

1.4.2 Exploratory Data Analysis

Let's use Matplotlib and Seaborn to explore the data.

```
[9]: insta_data.head(10)
```

```
[9]:
        Daily Time Spent on Instagram
                                                Daily Internet Usage \
                                           Age
                                   82.03
                                            41
                                                                187.53
     1
                                   80.03
                                            44
                                                                150.84
     2
                                   51.38
                                            59
                                                                158.56
     3
                                   77.07
                                            40
                                                                261.02
     4
                                   51.87
                                            50
                                                                119.65
     5
                                   45.99
                                            33
                                                                124.61
     6
                                   61.88
                                            45
                                                                108.18
     7
                                   68.01
                                            25
                                                                188.32
     8
                                   72.84
                                                                238.63
                                            26
     9
                                   73.84
                                            31
                                                                121.05
```

Ad feed Headline Location Sex \

| 0 | | attitude Pruittmout | | | 0 | | | | | | |
|---|---------------------------------------|----------------------|---------------|------------------|-------------|---|--|--|--|--|--|
| 1 | | Automated static | concept | Ch | ristinetown | 0 | | | | | |
| 2 | Object-ba | sed modular function | nalities | E | ast Anthony | 0 | | | | | |
| 3 | Face-t | o-face analyzing en | cryption | Ste | phenborough | 0 | | | | | |
| 4 | Tea | m-oriented dynamic | forecast | | Kevinberg | 0 | | | | | |
| 5 | Face- | to-face methodical | intranet | | South Mark | 1 | | | | | |
| 6 | Balanced dis | intermediate conglor | meration | North | Debrashire | 0 | | | | | |
| 7 | Amel | iorated actuating w | orkforce | | Kaylashire | 1 | | | | | |
| 8 | _ | ressive empowering | | | Wongland | 0 | | | | | |
| 9 | Progra | mmable uniform produ | uctivity | South | Jasminebury | 0 | | | | | |
| | Country | Timestamp (| Clicked o | n Feed | | | | | | | |
| 0 | · · · · · · · · · · · · · · · · · · · | 02-05-2016 07:00 | | 0 | | | | | | | |
| 1 | - | 23-07-2016 14:47 | | 1 | | | | | | | |
| 2 | Afghanistan | | | 1 | | | | | | | |
| 3 | Afghanistan 01-03-2016 10:01 0 | | | | | | | | | | |
| 4 | Afghanistan 10-03-2016 22:28 1 | | | | | | | | | | |
| 5 | <u> </u> | | | | | | | | | | |
| 6 | Afghanistan | 30-05-2016 08:02 | | 1 | | | | | | | |
| 7 | Afghanistan | | 0 | | | | | | | | |
| 8 | Albania | 08-07-2016 21:18 | | 0 | | | | | | | |
| 9 | Albania | 18-06-2016 17:23 | | 1 | | | | | | | |
| [10]· ir | usta data info | .() | | | | | | | | | |
| | [10]: insta_data.info() | | | | | | | | | | |
| <c< td=""><td>lass 'nandas d</td><td>core.frame.DataFrame</td><td>·'></td><td></td><td></td><td></td></c<> | lass 'nandas d | core.frame.DataFrame | ·'> | | | | | | | | |
| | - | entries, 0 to 999 | , | | | | | | | | |
| | • | otal 9 columns): | | | | | | | | | |
| # | | Non-Null Count Dtype | | | | | | | | | |
| | | | | | | | | | | | |
| 0 | Daily Time | 1000 nor | n-null | float64 | | | | | | | |
| 1 | - | | 1000 non-null | | int64 | | | | | | |
| 2 | - | 1000 non-null fi | | float64 | | | | | | | |
| 3 | • | 1000 non-null object | | object | | | | | | | |
| 4 | Location | 1000 non-null object | | • | | | | | | | |
| 5 | Sex | 1000 non-null int64 | | int64 | | | | | | | |
| 6 | Country | 1000 non-null object | | | | | | | | | |
| 7 | Timestamp | 1000 non-null object | | | | | | | | | |
| 8 | Clicked on | Feed | 1000 nor | 0 non-null int64 | | | | | | | |
| dtypes: float64(2), int64(3), object(4) | | | | | | | | | | | |
| memory usage: 70.4+ KB | | | | | | | | | | | |
| [11]: insta_data.describe() | | | | | | | | | | | |
| | | | | | | | | | | | |

1000.000000

36.009000

Age Daily Internet Usage $\$

1000.000000

180.000100

Daily Time Spent on Instagram

1000.000000

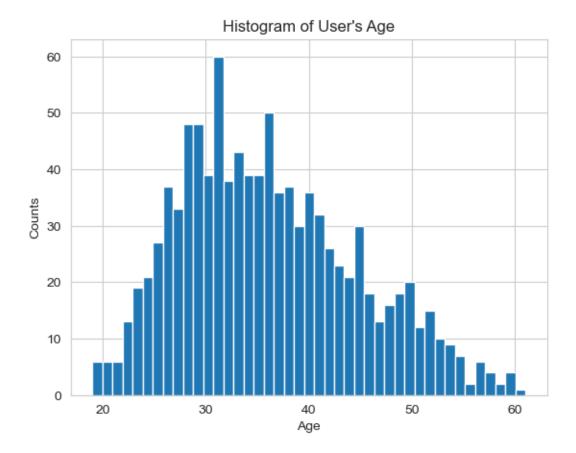
65.000200

[11]:

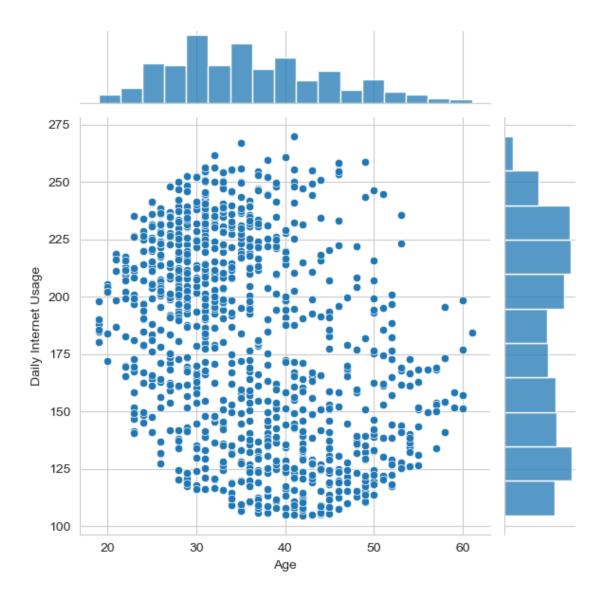
count

mean

```
std
                                  15.853615
                                                8.785562
                                                                     43.902339
     min
                                  32.600000
                                               19.000000
                                                                     104.780000
      25%
                                 51.360000
                                               29.000000
                                                                     138.830000
      50%
                                 68.215000
                                               35.000000
                                                                     183.130000
      75%
                                 78.547500
                                               42.000000
                                                                    218.792500
                                 91.430000
                                               61.000000
                                                                    269.960000
     max
                     Sex Clicked on Feed
      count 1000.000000
                               1000.00000
     mean
                0.481000
                                  0.50000
      std
                0.499889
                                  0.50025
     min
                0.000000
                                  0.00000
      25%
                0.000000
                                  0.00000
      50%
                0.000000
                                  0.50000
      75%
                1.000000
                                  1.00000
                1.000000
                                  1.00000
     max
[12]: # Let's build a histogram for the user's age
      sns.set_style('whitegrid')
      insta_data['Age'].hist(bins=43)
      plt.xlabel('Age')
      plt.ylabel('Counts')
      plt.title("Histogram of User's Age")
      plt.show()
```




[13]: ## Let's build a jointplot showing Daily Time Spent on Instagram versus Age.
sns.jointplot(x='Age',y='Daily Internet Usage',data=insta_data)
plt.show()

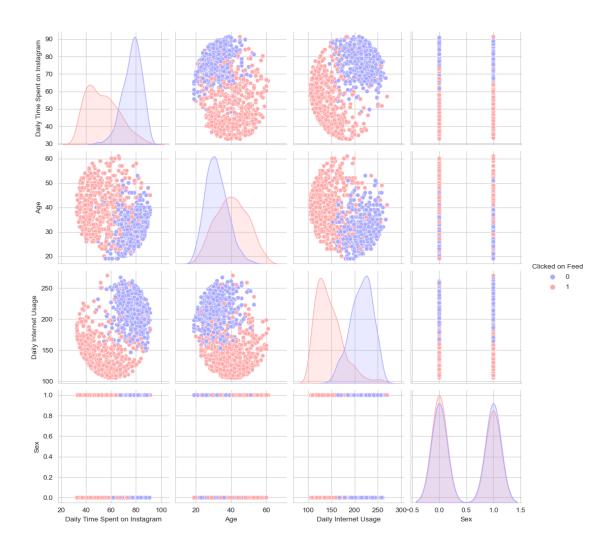


[14]: # Jointplot of 'Daily Time Spent on Site' vs. 'Daily Internet Usage'
sns.jointplot(x='Daily Time Spent on Instagram',y='Daily Internet

□ Usage',data=insta_data,color='green')
plt.show()



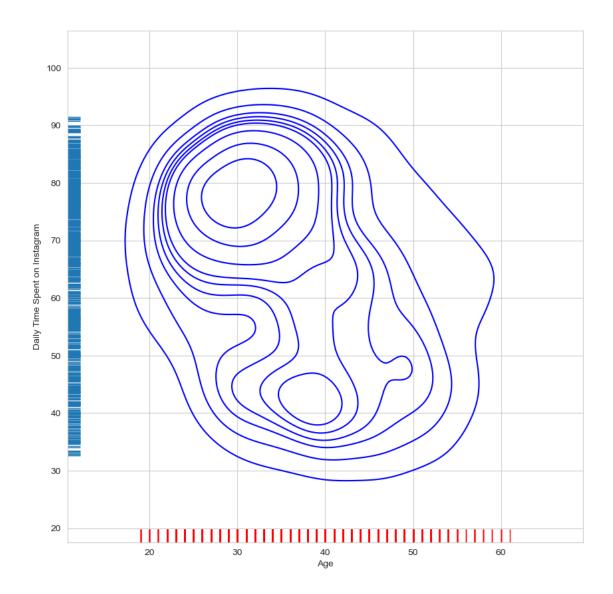
[15]: # Pairplot with a grouping defined by the 'Clicked on Ad' column feature.
sns.pairplot(insta_data,hue='Clicked on Feed',palette='bwr')
plt.show()



```
[16]: # How the user's age and the time spent on instagram are linked.

f, ax = plt.subplots(figsize=(10, 10))
sns.kdeplot(insta_data.Age, insta_data['Daily Time Spent on Instagram'],
color="b", ax=ax)
sns.rugplot(insta_data.Age, color="r", ax=ax)
sns.rugplot(insta_data['Daily Time Spent on Instagram'], vertical=True, ax=ax)
```

[16]: <AxesSubplot:xlabel='Age', ylabel='Daily Time Spent on Instagram'>



Data preparation

We have several categorical variables:

| [17]: | <pre>insta_data.select_dtypes(include='object').describe()</pre> | | | | | | | | | |
|-------|--|-----------|---------|----------|-----------|---------|------------------|--|--|--|
| [17]: | | | Ad feed | Headline | Location | Country | Timestamp | | | |
| | count | | | 1000 | 1000 | 1000 | 1000 | | | |
| | unique | | | 1000 | 969 | 237 | 997 | | | |
| | top | Intuitive | dynamic | attitude | Lisamouth | France | 26-05-2016 15:40 | | | |
| | freq | | | 1 | 3 | 9 | 2 | | | |

As we can see here, we have a lot of variability in the Ad feed headline, such as the location and country of the user, so for this example we will not use them.

You will eliminate these variables in question 4, and then in question 5 you will deal with the timestamp.

Return to top ## Question 4

Eliminate the variables Ad feed Headline, Location and Country from insta_data and save your resulting dataframe in a variable called df.

```
[18]: x = insta_data.select_dtypes(include='object').describe()
y = x.columns
y = y[:3]
y
```

[18]: Index(['Ad feed Headline', 'Location', 'Country'], dtype='object')

```
[19]: ### GRADED

### YOUR SOLUTION HERE

#insta_data.head(5)

df = insta_data
    df = df.drop(y, axis = 1)
    df.head(5)

###

### YOUR CODE HERE

###
```

```
Γ19]:
         Daily Time Spent on Instagram Age Daily Internet Usage
                                                                     Sex
                                  82.03
                                          41
                                                             187.53
                                                                       0
                                  80.03
      1
                                          44
                                                             150.84
                                                                       0
      2
                                  51.38
                                          59
                                                             158.56
                                                                       0
      3
                                  77.07
                                                             261.02
                                                                       0
                                          40
      4
                                  51.87
                                          50
                                                             119.65
                Timestamp Clicked on Feed
      0 02-05-2016 07:00
```

```
      0
      02-05-2016
      07:00
      0

      1
      23-07-2016
      14:47
      1

      2
      17-06-2016
      17:11
      1

      3
      01-03-2016
      10:01
      0

      4
      10-03-2016
      22:28
      1
```

Now let's see our new dataframe

```
[21]: df.head()
```

```
[21]:
         Daily Time Spent on Instagram
                                               Daily Internet Usage
                                          Age
                                                                      Sex
      0
                                   82.03
                                           41
                                                              187.53
                                                                         0
      1
                                   80.03
                                                              150.84
                                           44
                                                                         0
      2
                                   51.38
                                           59
                                                              158.56
                                                                         0
      3
                                   77.07
                                           40
                                                              261.02
                                                                         0
      4
                                   51.87
                                           50
                                                              119.65
                 Timestamp
                            Clicked on Feed
      0 02-05-2016 07:00
      1 23-07-2016 14:47
                                           1
      2 17-06-2016 17:11
                                           1
      3 01-03-2016 10:01
                                           0
      4 10-03-2016 22:28
                                           1
[22]: # Change the name of the variables so it's easier to work with them
      df.columns = ["time_instagram", "age", "time_internet", "sex", "timestamp", __

¬"clicked"]

[23]: df.head()
                               time_internet
[23]:
         time_instagram
                          age
                                                            timestamp
                                                                        clicked
                   82.03
                                       187.53
      0
                           41
                                                    02-05-2016 07:00
      1
                   80.03
                           44
                                       150.84
                                                    23-07-2016 14:47
                                                                              1
      2
                   51.38
                           59
                                       158.56
                                                    17-06-2016 17:11
                                                                              1
```

Return to top ## Question 5

77.07

51.87

3

Extract time features from the timestamp and create the variables:

261.02

119.65

• month: The month when the ad was seen

40

50

- day_of_month: The day of the month when the ad was seen
- day_of_week: The day of the week when the ad was seen
- hour: The hour of the day when the ad was seen

Extract the values as numeric features. Then add those features to the dataframe (df) as columns using the names indicated. Finally drop the timestamp column.

0

01-03-2016 10:01

10-03-2016 22:28

0

1

Hint: Transform timestamp to a datetime type object.

```
[24]: ### GRADED

### YOUR SOLUTION HERE
# Remember to cast 'timestamp' to a datetime object

df['timestamp'] = pd.to_datetime(df['timestamp'])
```

```
df['month'] = df['timestamp'].apply(lambda x:x.month)
      df['day_of_month'] = df['timestamp'].apply(lambda x:x.day)
      df["day_of_week"] = df['timestamp'].apply(lambda x:x.dayofweek)
      df['hour'] = df['timestamp'].apply(lambda x:x.hour)
      df = df.drop('timestamp', axis =1)
      ###
      ### YOUR CODE HERE
      ###
[25]: ###
      ### AUTOGRADER TEST - DO NOT REMOVE
[26]: df.head(10)
[26]:
                                                                       day_of_month
         time_instagram
                           age
                                time internet
                                                 sex
                                                      clicked
                                                                month
                   82.03
                            41
                                        187.53
                                                   0
                                                             0
                                                                     2
      1
                   80.03
                            44
                                        150.84
                                                   0
                                                             1
                                                                    7
                                                                                   23
      2
                   51.38
                            59
                                        158.56
                                                             1
                                                                    6
                                                                                   17
                                                   0
      3
                   77.07
                            40
                                        261.02
                                                   0
                                                             0
                                                                    1
                                                                                    3
      4
                                                                                    3
                   51.87
                            50
                                        119.65
                                                   0
                                                             1
                                                                   10
      5
                   45.99
                                                                    2
                                                                                   27
                            33
                                        124.61
                                                             1
                                                   1
      6
                   61.88
                            45
                                        108.18
                                                   0
                                                             1
                                                                    5
                                                                                   30
      7
                   68.01
                            25
                                        188.32
                                                   1
                                                             0
                                                                    1
                                                                                    1
      8
                   72.84
                            26
                                        238.63
                                                   0
                                                             0
                                                                    8
                                                                                    7
                            31
      9
                   73.84
                                        121.05
                                                             1
                                                                     6
                                                                                   18
         day_of_week
                       hour
                           7
      0
                    5
      1
                          14
      2
                    4
                          17
      3
                    6
                          10
      4
                    0
                          22
      5
                    5
                          15
      6
                    0
                           8
      7
                    4
                           3
      8
                    6
                          21
      9
                    5
                          17
```

You will now write Python code to split the data and fit a logistic regression model.

We will first need to split up our data into an X dataframe that contains the features to train on and a y Series with the target variable, which in this case will be the clicked column.

```
[27]: X = df[df.columns.difference(['clicked'])]
y = df['clicked']
```

```
[28]: print(X.head(5))
      print(y)
              day_of_month
                              day_of_week
                                             hour
                                                    month
                                                            sex
                                                                  time_instagram \
         age
                                                                            82.03
      0
          41
                                                 7
                                                        2
                                                              0
                                                        7
      1
          44
                          23
                                          5
                                                14
                                                              0
                                                                            80.03
      2
                                          4
                                                        6
                                                              0
                                                                            51.38
          59
                          17
                                                17
      3
          40
                           3
                                          6
                                                10
                                                        1
                                                              0
                                                                            77.07
                           3
      4
          50
                                          0
                                                22
                                                        10
                                                              0
                                                                            51.87
         time_internet
      0
                 187.53
      1
                 150.84
      2
                 158.56
      3
                 261.02
      4
                 119.65
      0
              0
      1
              1
      2
              1
      3
              0
      4
              1
      995
              1
      996
              1
      997
              1
      998
              0
      999
              0
      Name: clicked, Length: 1000, dtype: int64
```

1.4.3 Logistic Regression Model

Return to top ## Question 6

Split the data into a **training** set and a **testing** set. We will train our model on the training set and then use the test set to evaluate it. Use 30% of the data for testing and the remaining for training. Use 6 as a random state.

Return to top ## Question 7

Create and fit a Logistic Regression model using the training dataset. Save your fitted model in a variable called model.

Make sure to import LogisticRegression from sklearn. Use the lbfgs solver and random_state = 6 arguments when creating the model. Don't pass any other argument or changes to the logistic regression.

Return to top ## Question 8

Model evaluation

Create a predictions variable using your fitted model and the test dataset; call it y_pred. Then get the accuracy score of your predictions and save it in a variable called accuracy. Finally get the confusion matrix for your predictions and save it in a variable called confusion_mat.

Hint: Use the sklearn.metrics module.

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
### YOUR SOLUTION HERE

y_true = y_test
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_true, y_pred)
confusion_mat = confusion_matrix(y_true, y_pred)


###
```

```
### YOUR CODE HERE
###
```


```
[34]: print('Logistic regression accuracy: {}'.format(accuracy))
```

Logistic regression accuracy: 0.93

[[138 6] [15 141]]



Finally let's do a plot of our ROC while also getting the AUC.



[]: